

Engineering Databases

Lecture 4 – Foreign Key and Joins

November 9, 2022

M. Saeed Mafipour & Mansour Mehranfar

Contents of Lecture 3


- DML (partial):
 - update $\langle t1 \rangle$ set $\langle c1 \rangle = \langle v1 \rangle$ [, $\langle c2 \rangle = \langle c2 \rangle$, ...] [where $\langle \text{condition} \rangle$]
 - delete from $\langle t1 \rangle$ [where $\langle \text{condition} \rangle$]
 - insert into for multiple data sets
- Advanced ER-Mapping Schema
 - 1:N, N:1, and 1:1 relation can be eliminated
- The Relational Model by Codd
 - The formal groundings of SQL based on set theory and first-order predicate logic
 - Domains, attributes, relations and relational schemas
 - Projections: $\Pi_{\text{columns}}(\text{Relation})$, extract columns from a relation
 - Union: $R_i \cup R_j$, merge identical schemas
 - Selection: $\sigma_{\text{statement}}(\text{Relation})$, select sub-sets of a relation based on formulas

Foreign key integrity constraint

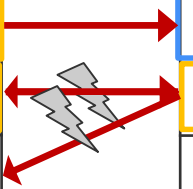
- An essential principle in relational design
- Needed to reflect real-world relations over multiple tables
- A foreign key is an attribute or attribute set
- It is identical to a primary key of another table
- The tables are linked via the foreign key constraint
- Is part of the DDL
- The DBMS ensures the data integrity in case of data changes

Foreign key integrity constraint

- Example:



| instructors | | | | offices | | |
|-------------|-------|------------|------|------------|------|-----------------|
| id | name | department | room | roomNumber | size | telephoneNumber |
| 1 | Bob | CMS | 10 | 10 | 10 | 22678 |
| 2 | Tom | CIE | 5 | 8 | 25 | 35131 |
| 3 | Alice | CMS | 8 | | | |



Foreign key integrity constraint

- Ensures that a value in a table is contained in some other table
- The syntax for **create table** with foreign key

```
CREATE TABLE <name of table> (  
    [other attribute definitions],  
    <name of foreign key> <type of foreign key> <integrity constraints>,  
    [other attributes that are part of the foreign key, comma separated]  
    FOREIGN KEY (<local foreign key attributes, comma separated>  
    REFERENCES <other table>  
    (<other table primary key attributes, comma separated>));
```

- The foreign key attributes (including type) have to match the other table's columns.
- The name of the attributes do not have to match!
- The other table must exist!

Foreign key integrity constraint

- Example:

↓

CREATE TABLE **offices** (

→ **roomNumber** INTEGER PRIMARY KEY,
size INTEGER NOT NULL,
telephonNumber VARCHAR(40) NOT NULL);

CREATE TABLE instructors (


id INTEGER PRIMARY KEY,
name VARCHAR(50) NOT NULL,
department VARCHAR(20),

→ **room** INTEGER,
FOREIGN KEY (**room**) **REFERENCES** **offices** (**roomNumber**));

| StudentTest | lec4f__offices |
|-------------|------------------------------|
| 🔑 | roomNumber : int(11) |
| # | size : int(11) |
| # | telephonNumber : varchar(40) |

| StudentTest | lec4f__instructors |
|-------------|--------------------------|
| 🔑 | id : int(11) |
| # | name : varchar(50) |
| # | department : varchar(20) |
| # | room : int(11) |

Foreign key integrity constraint exercise 1




- Create the tables `persons` and `hobbies` (moodle sql!)
- Create table `hasHobby` which includes two foreign key constraints but no primary key.
- Hints: It is not possible to have two attributes of the same name in a single table.
- This step is mandatory for later exercises, don't skip it!

Foreign key integrity constraint

- Ensures that a value in a table is contained in some other table
- The syntax for **alter table** with foreign key

```
ALTER TABLE <name of table> ADD  
    FOREIGN KEY (<local foreign key attributes, comma separated>  
    REFERENCES <other table>  
    (<other table primary key attributes, comma separated>);
```



Foreign key integrity constraint exercise 2



- Create a new table hobbyCategories
- Alter table hobbies:
 - Add the attribute typeld (new ER-Mapping rule!)
and the foreign key constraint on typeld for hobbyCategories id
 - Hint: Use a single alter table statement but multiple add statements
that are comma separated!
- This step is mandatory for later exercises, don't skip it!

Foreign key integrity constraint

- The foreign key ensures data integrity



- Examples:

delete from offices where roomNumber=5

update instructors set room=10 where department='CMS'

insert into instructors values (4, Mark, 'CMS', 10)


- In all cases, an error is given.

Foreign key integrity constraint

- Control updates and changes of other tables
- The syntax for **update and delete** actions in the foreign key definition
FOREIGN KEY (<local foreign>)
REFERENCES <other table> (<other table primary key attributes>)
[ON DELETE [RESTRICT | CASCADE | SET NULL | NO ACTION | SET DEFAULT]]
[ON UPDATE [RESTRICT | CASCADE | SET NULL | NO ACTION | SET DEFAULT]]
- Example:
foreign key (room) references offices (roomNumber) on delete cascade
- Restrict is default, means not allowed
- Cascade means that the other table's updates are propagated
- Set Null will set the value to null (if no NOT NULL constraint is given)
- No action will change nothing
- Set default will set the value to a constraint defined default value

Foreign key integrity constraint

- The foreign key ensures data integrity



- Example for delete from offices where roomNumber=5
 - Restrict: error
 - Cascade: Bob und Tom are deleted!
 - Set null: room in Bob and Tom are null (if no NOT NULL constrain is given)
 - No action: the room in Bob and Tom is 5
 - Set default: the room in Bob and Tom is set to a predefined default value

Relational Algebra

- Formal definitions of

Cartesian product \times

Renaming ρ

θ -Join \bowtie_{θ} and equi-Join $\bowtie_{\theta \text{ is } =}$

(Natural) Join \bowtie

Right \ltimes and left semi join \ltimes

Left \Join , right \Join , and full outer join \Join

Cartesian Product (Cross Product)

In general: $R \times S$

- Contains $|R| \cdot |S|$ combinations (all) of tuples from R and S.
- The resulting relation $R \times S$ is the combination of their attributes
- The SQL syntax for the Cartesian product:
 $\langle \text{Relation 1} \rangle \text{ CROSS JOIN } \langle \text{Relation 2} \rangle$
 or
 $\langle \text{Relation 1} \rangle \text{ JOIN } \langle \text{Relation 2} \rangle$
- CROSS JOIN and JOIN are synonyms in MySQL (but not in other DBMS)

Cartesian Product (Cross Product)

- Example:

instructors \times offices


SELECT * FROM instructors CROSS JOIN offices

SELECT * FROM instructors JOIN offices

| instructors | |
|-------------|-------|
| name | chair |
| Bob | CMS |
| Tom | CIE |



| offices | |
|---------|------|
| room | size |
| 005 | 14 |
| 008 | 25 |



| instructors \times offices | | | |
|------------------------------|-------|------|------|
| name | chair | room | size |
| Bob | CMS | 005 | 14 |
| Tom | CIE | 005 | 14 |
| Bob | CMS | 008 | 25 |
| Tom | CIE | 008 | 25 |

- Multiple cross products: persons \times hasHobby \times hobbies

Renaming

- Problem occurs when using the same relation twice.
- E.g. the instructors \times instructors will fail

~~SELECT * FROM instructors JOIN instructors~~

In general: $\rho_{new\ name}$ (Relation)

- The SQL syntax for the renaming is:
 <Relation> AS <New Name>

- Example:

$\rho_{inst1}(Instructors) \times \rho_{inst2}(Instructors)$

SELECT * FROM instructors AS inst1 JOIN instructors AS inst2

- We already used the ρ operation to give unnamed relation a name (nesting!)

θ -Join and Equi-Join

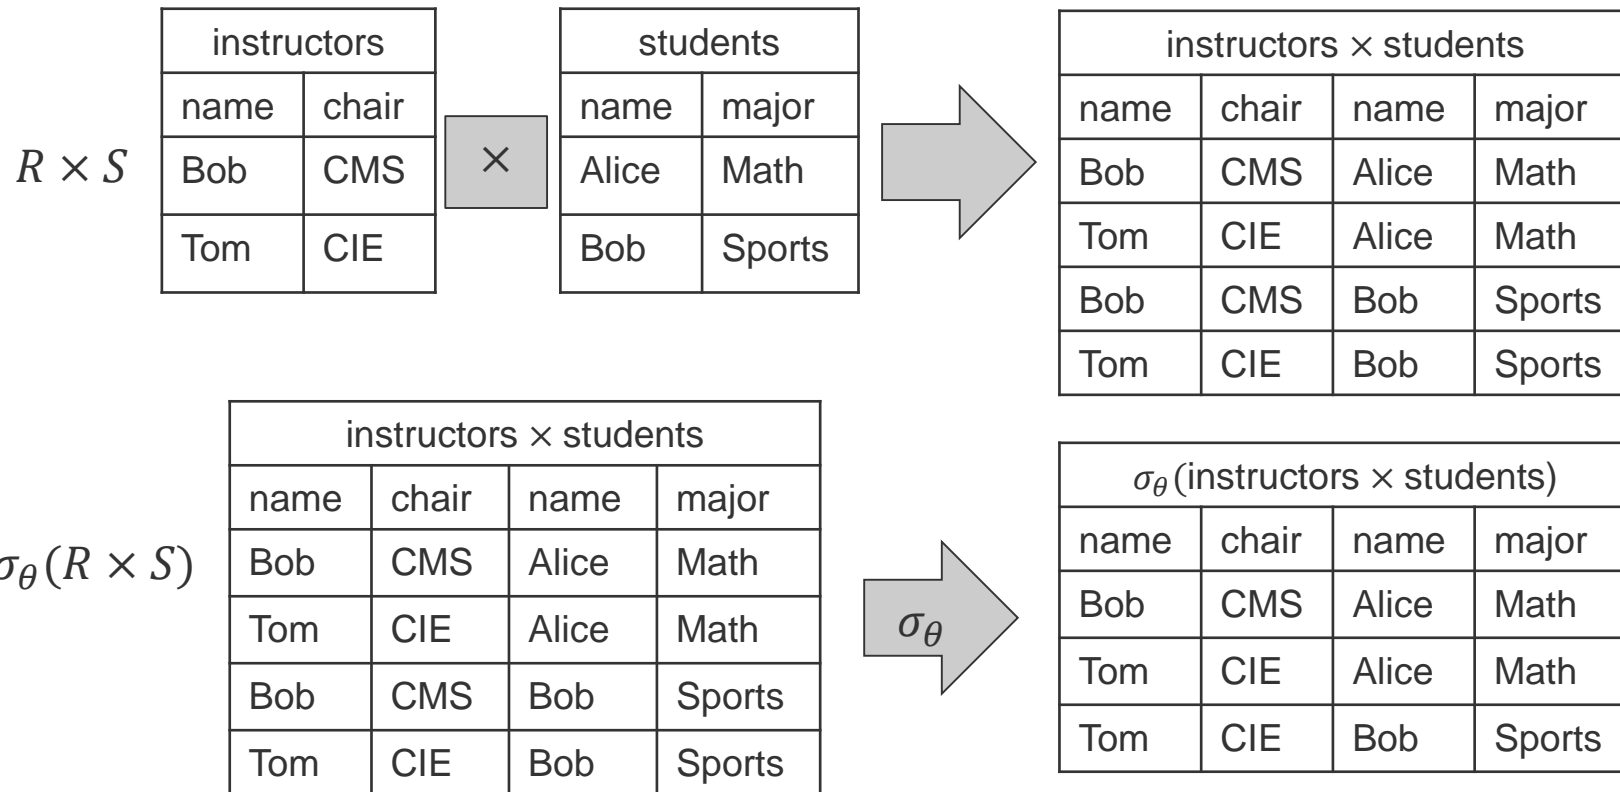
In general: $R \bowtie_{\theta} S$ is a θ -Join. Also: $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$

- Not all rows of the Cartesian product are desired
- Result contains all entities of $R \times S$ for which formula θ is fulfilled
- In case the condition comprises only = it is an equi-join
- The SQL syntax for the Cartesian product:
 SELECT * FROM <relation1> JOIN < relation2 >
 ON <condition1> [AND|OR more Conditions]
 or
 SELECT * FROM <table1> INNER JOIN <table2>
 ON <condition1> [AND|OR more Conditions]
- INNER JOIN and JOIN are synonyms in MySQL (but not in other DBMS)

θ -Join

- Example: $\text{instructors} \bowtie_{\text{instructors.name} \neq \text{students.name}} \text{students}$

SELECT * FROM instructors JOIN students ON instructors.name != students.name



Equi-Join

- Example: $\text{instructors} \bowtie_{\text{instructors.name=students.name}} \text{students}$


SELECT * FROM instructors JOIN students ON instructors.name = students.name

or

SELECT * FROM instructors INNER JOIN students ON
instructors.name = students.name


or

SELECT * FROM instructors CROSS JOIN students ON
instructors.name = students.name



- Including renaming: select * from instructors as i join students s on i.name = s.name

Equi-Join exercise



- If not already done, create tables persons, hasHobby, hobbies, and hobbyCategories.
- Populate the tables with data (provided in the lecture)

- Execute the SQL commands:

$$\text{hobbies} \bowtie_{id = hobbyId} \text{hasHobby}$$

$$\rho_h(\text{hobbies}) \bowtie_{typeId = id} \text{hobbyCategories}$$

$$\sigma_{age > 30}(\text{persons} \bowtie_{id = personId} \text{hasHobby})$$

$$\Pi_{age, name}(\text{persons} \bowtie_{id = personId} \text{hasHobby})$$

$$\text{hobbies} \bowtie_{id = hobbyId} \text{hasHobby} \bowtie_{personId = id} \text{persons}$$

$$\rho_h(\text{hobbies}) \bowtie_{id = hobbyId} \rho_{hH}(\text{hasHobby}) \bowtie_{personId = id} \rho_p(\text{persons})$$

Natural Join

In general: $R \bowtie S$

- Same as equi-join with columns, having the same name, set equal.
- Is associative and communicative: $A \bowtie B \bowtie C = A \bowtie (B \bowtie C) = B \bowtie (A \bowtie C)$
- The SQL syntax for the natural join

`SELECT * FROM <relation1> NATURAL JOIN <relation2 >`

- Example:

`instructors \bowtie students`

`SELECT * FROM instructors NATURAL JOIN students`

- Exercise:

`persons \bowtie hobbies (this joins on name and id!)`


Semi Join

In general: left semi join $R \bowtie S = \Pi_R(R \bowtie S)$, right semi join: $R \bowtie S = \Pi_S(R \bowtie S)$

- Selects one table's columns that were matched during a natural join.
- The SQL syntax for the left semi join
SELECT <r1>.* FROM <r1> NATURAL JOIN <r2>
- The SQL syntax for the right semi join
SELECT <r2>.* FROM <r1> NATURAL JOIN <r2>
- Example:
instructors \bowtie students (left semi join)
instructors \bowtie students (right semi join)

| instructors | |
|-------------|-------|
| name | chair |
| Bob | CMS |
| Tom | CIE |

| students | |
|----------|--------|
| name | major |
| Alice | Math |
| Bob | Sports |



| instructors \bowtie students | |
|--------------------------------|-------|
| name | chair |
| Bob | CMS |

| instructors \bowtie students | |
|--------------------------------|--------|
| name | major |
| Bob | Sports |

Outer Join

In general: left outer join $\bowtie\leftarrow$, right outer join $\rightarrow\bowtie$, and full outer join \bowtie

- Compared to Equi-Join, outer joins keep the unmatched entries
- The entries are taken from the left (left outer join), the right (right outer join) or both (full outer join) relations.
- The SQL syntax for the outer joins:
 - left outer join: `SELECT * FROM <r1> LEFT JOIN <r2> ON <Condition>`
 - right outer join: `SELECT * FROM <r1> RIGHT JOIN <r2> ON <Condition>`
 - full outer join: `<left outer join> UNION [ALL] <right outer join>`
- Adding ALL after UNION will keep identical rows

Outer Join

- Example: (on name)
 - instructors \bowtie students (right outer join)
 - instructors \Join students (left outer join)
 - instructors \Join students (full outer join)

| instructors | |
|-------------|-------|
| name | chair |
| Bob | CMS |
| Tom | CIE |

| students | |
|----------|--------|
| name | major |
| Alice | Math |
| Bob | Sports |

| instructors \Join students | | | |
|------------------------------|-------|-------|--------|
| name | chair | name | major |
| Bob | CMS | Bob | Sports |
| NULL | NULL | Alice | Math |

| instructors \Join students | | | |
|------------------------------|-------|------|--------|
| name | chair | name | major |
| Bob | CMS | Bob | Sports |
| Tom | CIE | NULL | NULL |

| Instructors \Join students | | | |
|------------------------------|-------|-------|--------|
| name | chair | name | major |
| Bob | CMS | Bob | Sports |
| Tom | CIE | NULL | NULL |
| Bob | CMS | Bob | Sports |
| NULL | NULL | Alice | Math |

Homework (solution is on moodle!)

- There are 4 relations: Students, Professors, Rooms, Classes
- Students have a name and a studentId.
- Professors have a name, faculty and a profId
- Rooms have a name, a number of seats, and a roomId.
- Classes have a name, a limit on the number of students, a room, a professor, and a classId.
 - Store **ONLY** the ids of the professors and rooms in this table. No names,...
 - Room and Professor are foreign keys, so only valid ids can be stored.
- Draw an entity relationship diagram.
- Create the tables
- Add 3 Students, 3 Professors, 3 Rooms and 3 Classes
- Output all professors who teach a course.
- Output all courses taught by one of the professors you entered.
- Output all courses with room names and professor names.
- Output one list containing the names of all professors and students

*Also write all
queries in
relational
algebra notation*



End of Lecture

Thank you for your attention